

Here's a complete implementation of a **Variational Autoencoder (VAE)** for Fashion-MNIST reconstruction in PyTorch, with visualization:

```
# File: 03_VAE_FashionMNIST_Reconstruction.py
# Topic: VAE implementation for Fashion-MNIST reconstruction using
PyTorch

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
from torchvision import transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
import numpy as np
from torch.utils.tensorboard import SummaryWriter
import os

# Set random seed for reproducibility
torch.manual_seed(42)
np.random.seed(42)

# Device configuration
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

# Hyperparameters
latent_dim = 20
batch_size = 128
epochs = 30
learning_rate = 1e-3
image_size = 28 * 28 # Fashion-MNIST dimensions

# Create output directory
os.makedirs('vae_results', exist_ok=True)
os.makedirs('vae_results/reconstructions', exist_ok=True)
writer = SummaryWriter('vae_results/runs')

# Fashion-MNIST dataset
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,)) # Scale to [-1, 1]
])

train_dataset = torchvision.datasets.FashionMNIST(
    root='./data', train=True, transform=transform, download=True)
train_loader = DataLoader(train_dataset, batch_size=batch_size,
    shuffle=True)

# VAE Model
class VAE(nn.Module):
```

```

def __init__(self):
    super(VAE, self).__init__()

    # Encoder
    self.fc1 = nn.Linear(image_size, 400)
    self.fc21 = nn.Linear(400, latent_dim) #  $\mu$  (mean)
    self.fc22 = nn.Linear(400, latent_dim) # logvar

    # Decoder
    self.fc3 = nn.Linear(latent_dim, 400)
    self.fc4 = nn.Linear(400, image_size)

def encode(self, x):
    h1 = F.relu(self.fc1(x))
    return self.fc21(h1), self.fc22(h1)

def reparameterize(self, mu, logvar):
    std = torch.exp(0.5 * logvar)
    eps = torch.randn_like(std)
    return mu + eps * std

def decode(self, z):
    h3 = F.relu(self.fc3(z))
    return torch.tanh(self.fc4(h3)) # Tanh to match input range [-1,1]

def forward(self, x):
    mu, logvar = self.encode(x.view(-1, image_size))
    z = self.reparameterize(mu, logvar)
    return self.decode(z), mu, logvar

# Loss function
def loss_function(recon_x, x, mu, logvar):
    BCE = F.mse_loss(recon_x, x.view(-1, image_size), reduction='sum')
    KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
    return BCE + KLD

# Initialize model and optimizer
model = VAE().to(device)
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

# Fixed input for visualization
fixed_input = next(iter(train_loader))[0][:32].to(device)

def save_reconstructions(epoch, original, reconstructed):
    """Save original and reconstructed images for comparison"""
    original = original.cpu().numpy()
    reconstructed = reconstructed.cpu().numpy()

    plt.figure(figsize=(10, 4))
    for i in range(8): # Display first 8 samples
        # Original images
        ax = plt.subplot(2, 8, i + 1)

```

```

plt.imshow(original[i].reshape(28, 28), cmap='gray')
plt.axis('off')

# Reconstructed images
ax = plt.subplot(2, 8, i + 8 + 1)
plt.imshow(reconstructed[i].reshape(28, 28), cmap='gray')
plt.axis('off')

plt.savefig(f'vae_results/reconstructions/epoch_{epoch:03d}.png')
plt.close()

# Training loop
for epoch in range(epochs):
    model.train()
    train_loss = 0
    for batch_idx, (data, _) in enumerate(train_loader):
        data = data.to(device)

        optimizer.zero_grad()
        recon_batch, mu, logvar = model(data)
        loss = loss_function(recon_batch, data, mu, logvar)
        loss.backward()
        train_loss += loss.item()
        optimizer.step()

        if batch_idx % 100 == 0:
            print(f'Train Epoch: {epoch} [{batch_idx *
len(data)}/{len(train_loader.dataset)} '
                  f'({100. * batch_idx /
len(train_loader):.0f}%)]\tLoss: {loss.item() / len(data):.6f}')

# Validation and visualization
model.eval()
with torch.no_grad():
    # Calculate validation loss
    val_loss = 0
    for data, _ in train_loader: # Using same data for simplicity
        data = data.to(device)
        recon, mu, logvar = model(data)
        val_loss += loss_function(recon, data, mu, logvar).item()

# Generate reconstructions
recon_fixed, _, _ = model(fixed_input)

# Save reconstructions
save_reconstructions(epoch, fixed_input, recon_fixed)

# TensorBoard logging
writer.add_scalar('Training Loss', train_loss /
len(train_loader.dataset), epoch)
writer.add_scalar('Validation Loss', val_loss /
len(train_loader.dataset), epoch)

```

```

# Save sample to TensorBoard
comparison = torch.cat([
    fixed_input[:8],
    recon_fixed.view(-1, 1, 28, 28)[:8]
])
writer.add_images(
    'Original vs Reconstructed',
    (comparison + 1) / 2, # Scale to [0,1]
    epoch
)

print(f'====> Epoch: {epoch} Average loss: {train_loss /
len(train_loader.dataset):.4f}')

# Save model
torch.save(model.state_dict(), 'vae_results/vae_fashionmnist.pth')
writer.close()

# Generate new samples
with torch.no_grad():
    sample = torch.randn(64, latent_dim).to(device)
    generated = model.decode(sample).cpu()

plt.figure(figsize=(10, 10))
for i in range(64):
    plt.subplot(8, 8, i + 1)
    plt.imshow(generated[i].view(28, 28), cmap='gray')
    plt.axis('off')
plt.savefig('vae_results/generated_samples.png')
plt.close()

```

Key Components Explained:

1. VAE Architecture:

- **Encoder:** Maps input to latent space (μ and $\log\sigma$)
- **Reparameterization Trick:** Enables backpropagation through random sampling
- **Decoder:** Reconstructs input from latent space

2. Loss Function:

- **Reconstruction Loss (MSE):** Measures pixel-wise difference
- **KL Divergence:** Regularizes latent space to be close to standard normal

3. Fashion-MNIST Specifics:

- Input images normalized to $[-1, 1]$ range
- Tanh activation in decoder matches input range
- 28x28 grayscale images

4. Visualization Features:

- Saves original vs reconstructed comparisons each epoch
- Generates new samples from random latent vectors
- TensorBoard integration for loss tracking

Expected Output:

```
vae_results/
├── reconstructions/
│   ├── epoch_000.png
│   ├── epoch_001.png
│   └── ...
├── runs/ (TensorBoard logs)
├── generated_samples.png
└── vae_fashionmnist.pth
```

Monitoring Training:

```
tensorboard --logdir vae_results/runs
```

Common Improvements:

1. For Better Quality:

```
# Replace fc layers with convolutional layers
class ConvVAE(nn.Module):
    def __init__(self):
        super().__init__()
        # Encoder
        self.encoder = nn.Sequential(
            nn.Conv2d(1, 32, 3, stride=2, padding=1),
            nn.ReLU(),
            nn.Conv2d(32, 64, 3, stride=2, padding=1),
            nn.ReLU()
        )
        # Latent space
        self.fc_mu = nn.Linear(64*7*7, latent_dim)
        self.fc_logvar = nn.Linear(64*7*7, latent_dim)
        # Decoder
        self.decoder_input = nn.Linear(latent_dim, 64*7*7)
        self.decoder = nn.Sequential(
            nn.ConvTranspose2d(64, 32, 3, stride=2, padding=1,
                               output_padding=1),
            nn.ReLU(),
            nn.ConvTranspose2d(32, 1, 3, stride=2, padding=1,
                               output_padding=1),
```

```
nn.Tanh( )  
)
```

2. For Different Objectives:

- Use BCE loss instead of MSE for sharper images
- Add perceptual loss with a pre-trained network

3. For Advanced Use:

- Implement β -VAE for disentangled representations
- Add classifier guidance for conditional generation

Would you like me to modify any aspect of this implementation or add specific features like:

- Latent space interpolation visualization?
- Anomaly detection capabilities?
- Quantitative evaluation metrics (FID, etc.)?