

Multi-Agent Environment Simulation with RLlib (PyTorch)

Below is a comprehensive example of setting up a multi-agent environment with RLlib for both cooperative and competitive tasks using PyTorch.

1. Setting Up a Custom Multi-Agent Environment

First, let's create a custom environment that supports both cooperative and competitive scenarios:

```
import numpy as np
import gym
from gym.spaces import Discrete, Box
from ray.rllib.env.multi_agent_env import MultiAgentEnv

class MultiAgentGridWorld(MultiAgentEnv):
    """
    A grid world where agents can cooperate or compete based on reward
    structure.
    """
    def __init__(self, config=None):
        config = config or {}
        self.size = config.get("grid_size", 5)
        self.max_steps = config.get("max_steps", 100)
        self.cooperative = config.get("cooperative", True)
        self.num_agents = config.get("num_agents", 2)

        # Define observation and action spaces
        self.observation_space = Box(low=0, high=1, shape=(
            self.size*self.size + 2*self.num_agents,))
        self.action_space = Discrete(4) # Up, Down, Left, Right

        # Agents IDs
        self.agents = [f"agent_{i}" for i in range(self.num_agents)]

        # Environment state
        self.reset()

    def reset(self):
        self.steps = 0
        self.agent_positions = {
            agent: np.random.randint(0, self.size, size=2)
            for agent in self.agents
        }

        # Place targets
        if self.cooperative:
            # Shared target for cooperation
```

```

        self.target = np.random.randint(0, self.size, size=2)
    else:
        # Separate targets for competition
        self.targets = {
            agent: np.random.randint(0, self.size, size=2)
            for agent in self.agents
        }

    return self._get_obs()

def _get_obs(self):
    obs = {}
    grid = np.zeros((self.size, self.size))

    for agent in self.agents:
        # Agent's own position
        pos = self.agent_positions[agent]
        grid[pos[0], pos[1]] = 1

        # Other agents' positions
        other_agents = [a for a in self.agents if a != agent]
        for i, other in enumerate(other_agents):
            opos = self.agent_positions[other]
            grid[opos[0], opos[1]] = 0.5

        # Target information
        if self.cooperative:
            grid[self.target[0], self.target[1]] = -1
            target_info = self.target
        else:
            grid[self.targets[agent][0], self.targets[agent][1]] =
-1
            target_info = self.targets[agent]

        # Flatten grid and add agent-specific info
        agent_obs = np.concatenate([
            grid.flatten(),
            pos,
            target_info
        ])
        obs[agent] = agent_obs

    return obs

def step(self, actions):
    rewards = {agent: 0 for agent in self.agents}
    dones = {"__all__": False}
    self.steps += 1

    # Move agents
    for agent, action in actions.items():
        pos = self.agent_positions[agent]

```

```

# Action effects
if action == 0: # Up
    pos[0] = max(0, pos[0] - 1)
elif action == 1: # Down
    pos[0] = min(self.size - 1, pos[0] + 1)
elif action == 2: # Left
    pos[1] = max(0, pos[1] - 1)
elif action == 3: # Right
    pos[1] = min(self.size - 1, pos[1] + 1)

# Check if reached target
if self.cooperative:
    if np.array_equal(pos, self.target):
        rewards[agent] = 1
        if all(np.array_equal(self.agent_positions[a],
self.target) for a in self.agents):
            rewards = {a: 10 for a in self.agents} # Big
reward for all if all reach target
            done["__all__"] = True
else:
    if np.array_equal(pos, self.targets[agent]):
        rewards[agent] = 10
        # Negative reward for others in competitive mode
        for other in self.agents:
            if other != agent:
                rewards[other] = -5
        done["__all__"] = True

# Timeout
if self.steps >= self.max_steps:
    done["__all__"] = True

obs = self._get_obs()
info = {} # Additional info if needed

return obs, rewards, done, info

```

PROF

2. Configuring and Training with RLlib

Now let's set up the RLlib training configuration for both cooperative and competitive scenarios:

```

import ray
from ray import tune
from ray.rllib.agents.ppo import PPOTrainer
from ray.rllib.models import ModelCatalog
from ray.rllib.models.torch.torch_modelv2 import TorchModelV2
import torch
import torch.nn as nn

# Define a custom neural network model

```

```

class CustomModel(TorchModelV2, nn.Module):
    def __init__(self, obs_space, action_space, num_outputs,
model_config, name):
        TorchModelV2.__init__(self, obs_space, action_space,
num_outputs, model_config, name)
        nn.Module.__init__(self)

        self.fcnet = nn.Sequential(
            nn.Linear(obs_space.shape[0], 64),
            nn.ReLU(),
            nn.Linear(64, 64),
            nn.ReLU(),
        )

        self.action_out = nn.Linear(64, num_outputs)
        self.value_out = nn.Linear(64, 1)

    def forward(self, input_dict, state, seq_lens):
        features = self.fcnet(input_dict["obs"])
        self._value_out = self.value_out(features)
        return self.action_out(features), state

    def value_function(self):
        return self._value_out.squeeze(1)

# Register the custom model
ModelCatalog.register_custom_model("custom_model", CustomModel)

def train_multi_agent(config, coop=True):
    # Initialize Ray
    ray.init(ignore_reinit_error=True)

    # Configuration
    config = {
        "env": MultiAgentGridWorld,
        "env_config": {
            "grid_size": 5,
            "num_agents": 2,
            "cooperative": coop,
            "max_steps": 100,
        },
        "multiagent": {
            "policies": {
                # Define one policy per agent (could share or have
                # separate policies)
                f"policy_{i}": (
                    None, # Use default obs/act spaces from env
                    Box(0, 1, (5*5 + 2*2,)), # Custom obs space
                    Discrete(4), # Action space
                    {
                        "model": {
                            "custom_model": "custom_model",
                        },
                    },
                )
            }
        }
    }

```

```

        "gamma": 0.95,
    }
    ) for i in range(2)
},
    "policy_mapping_fn": lambda agent_id:
f"policy_{int(agent_id.split('_')[1])}",
    },
    "framework": "torch",
    "num_workers": 3,
    "num_envs_per_worker": 5,
    "train_batch_size": 4000,
    "rollout_fragment_length": 200,
    "sgd_minibatch_size": 256,
    "lr": 1e-4,
}

# Select algorithm (PPO in this case)
trainer = PPO TorchTrainer(config=config)

# Training loop
for i in range(10): # 10 training iterations
    result = trainer.train()
    print(f"Iteration {i}:")
    print(f" - Total reward: {result['episode_reward_mean']}")
    print(f" - Episode length: {result['episode_len_mean']}")

    # Optionally save the model
    if i % 5 == 0:
        checkpoint = trainer.save()
        print(f"Checkpoint saved at {checkpoint}")

# Shutdown Ray when done
ray.shutdown()

# Train cooperative agents
print("Training cooperative agents...")
train_multi_agent(coop=True)

# Train competitive agents
print("\nTraining competitive agents...")
train_multi_agent(coop=False)

```

PROF

3. Evaluating the Trained Agents

After training, you can evaluate the agents:

```

def evaluate_agents(coop=True):
    # Initialize Ray
    ray.init(ignore_reinit_error=True)

```

```

# Create environment
env = MultiAgentGridWorld({
    "grid_size": 5,
    "num_agents": 2,
    "cooperative": coop,
    "max_steps": 100,
})

# Load trained policies
if coop:
    checkpoint_path = "path_to_cooperative_checkpoint"
else:
    checkpoint_path = "path_to_competitive_checkpoint"

trainer = PPOTorchTrainer(config={
    "env": MultiAgentGridWorld,
    "framework": "torch",
    "num_workers": 0,
})
trainer.restore(checkpoint_path)

# Run evaluation episodes
for ep in range(3): # Run 3 evaluation episodes
    obs = env.reset()
    done = {"__all__": False}
    total_rewards = {agent: 0 for agent in env.agents}

    while not done["__all__"]:
        actions = {}
        for agent_id in obs.keys():
            policy_id = f"policy_{int(agent_id.split('_')[1])}"
            actions[agent_id] =
trainer.compute_action(obs[agent_id], policy_id=policy_id)

        obs, rewards, done, _ = env.step(actions)

        for agent_id in rewards:
            total_rewards[agent_id] += rewards[agent_id]

    print(f"Episode {ep + 1} rewards:")
    for agent_id, reward in total_rewards.items():
        print(f"  {agent_id}: {reward}")

    ray.shutdown()

# Evaluate cooperative agents
print("Evaluating cooperative agents...")
evaluate_agents(coop=True)

# Evaluate competitive agents
print("\nEvaluating competitive agents...")
evaluate_agents(coop=False)

```

Key Concepts in This Implementation:

1. **Multi-Agent Environment:** The custom environment supports both cooperative and competitive scenarios through a config flag.
2. **Policy Sharing:** Agents can share policies or have individual policies based on the configuration.
3. **Custom Neural Network:** A PyTorch model is defined for policy approximation.
4. **Training Configuration:** RLlib's PPO algorithm is used with appropriate multi-agent settings.
5. **Evaluation:** Separate evaluation function to test trained agents.

This setup can be extended with:

- More complex environment dynamics
- Additional agents
- Different reward structures
- Hierarchical or communication-based policies
- Parameter sharing between agents

Remember to adjust hyperparameters (learning rate, batch sizes, etc.) based on your specific requirements and environment complexity.