Managing Applications in Kubernetes with Deployments and ReplicaSets

Kubernetes Deployments provide a declarative way to manage applications by controlling ReplicaSets, which in turn manage Pods. Here's how it works and a YAML example:

Key Concepts

- 1. **Deployment**: Manages the desired state of your application
- 2. **ReplicaSet**: Ensures the specified number of pod replicas are running
- 3. Rolling Updates: Deployments support zero-downtime updates
- 4. Rollbacks: Ability to revert to previous versions

YAML Example

Here's a complete example Deployment YAML for a simple web application:

```
apiVersion: apps/v1
kind: Deployment
metadata:
 name: webapp-deployment
 labels:
   app: webapp
spec:
 replicas: 3
 selector:
   matchLabels:
     app: webapp
 strategy:
   type: RollingUpdate
   rollingUpdate:
     maxSurge: 1
     maxUnavailable: 0
  template:
   metadata:
     labels:
       app: webapp
   spec:
     containers:
      - name: webapp-container
       image: nginx:1.19.1
       ports:
       - containerPort: 80
       resources:
          requests:
           cpu: "100m"
           memory: "128Mi"
           cpu: "200m"
            memory: "256Mi"
```

Step-by-Step Management

1. Create the Deployment

```
kubectl apply -f deployment.yaml
```

2. Check Deployment Status

```
kubectl get deployments
kubectl describe deployment webapp-deployment
```

3. View ReplicaSets

```
kubectl get replicasets
kubectl describe replicaset <replicaset-name>
```

4. View Pods

```
kubectl get pods --selector app=webapp
```

5. Update the Deployment (Rolling Update)

Edit the deployment.yaml file to change the image version:

```
image: nginx:1.20.0
```

Then apply:

```
kubectl apply -f deployment.yaml
```

6. Check Rollout Status

kubectl rollout status deployment/webapp-deployment

7. Rollback if Needed

kubectl rollout undo deployment/webapp-deployment

PROF

8. Scale the Deployment

```
kubectl scale deployment webapp-deployment --replicas=5
```

Key Benefits

- 1. **Self-healing**: Automatically replaces failed pods
- 2. **Scaling**: Easily scale up/down with simple commands
- 3. **Updates**: Controlled rollout of new versions
- 4. Rollbacks: Quickly revert problematic updates

Kubernetes Ingress: Managing External Access to Services

Kubernetes Ingress is an API object that provides HTTP/HTTPS routing rules to manage external access to services in a cluster. It acts as a smart traffic router, offering a more flexible alternative to NodePort or LoadBalancer services.

How Ingress Works

- 1. Ingress Resource: You define rules in a YAML manifest that specify how to route incoming requests
- 2. Ingress Controller: A pod that watches for Ingress resources and configures the underlying load balancer/proxy
- 3. External Access: Clients connect to the Ingress endpoint which then routes to appropriate services

Key Components

1. Ingress Resource (YAML Definition)

Defines the routing rules, TLS certificates, and backend services.

2. Ingress Controller

Actual implementation that processes the rules (common options: NGINX, Traefik, HAProxy, ALB, Istio)

3. Backend Services

Your regular ClusterIP services that Ingress routes to

Example Ingress YAML

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
 name: webapp-ingress
 annotations:
   nginx.ingress.kubernetes.io/rewrite-target: /
spec:
 tls:
  - hosts:
   - myapp.example.com
   secretName: tls-secret
  - host: myapp.example.com
   http:
     paths:
      - path: /webapp
       pathType: Prefix
       backend:
         service:
           name: webapp-service
            port:
             number: 80
      - path: /api
       pathType: Prefix
       backend:
          service:
           name: api-service
            port:
             number: 8080
```

How Ingress Helps Manage External Access

1. Host-Based Routing: Route traffic based on domain name

```
    host: api.example.com → api-service
    host: app.example.com → webapp-service
```

2. Path-Based Routing: Route based on URL path

```
- path: /v1 → legacy-service
- path: /v2 → modern-service
```

3. TLS Termination: Handle SSL certificates at the ingress level

```
tls:
- hosts: [myapp.example.com]
secretName: tls-secret
```

- 4. Load Balancing: Distribute traffic across service pods
- 5. Name-Based Virtual Hosting: Serve multiple websites from single IP
- 6. Traffic Control: Annotations for rate limiting, redirects, rewrites

Benefits Over Other Approaches

Feature	Ingress	LoadBalancer	NodePort
Single IP	\mathscr{O}	× (per service)	×
Path routing	\mathscr{C}	×	×
Host routing	4	×	×
TLS termination	V	V	×
Cost efficiency	V	×	\mathscr{O}

Typical Workflow

- 1. Deploy an Ingress Controller (once per cluster)
- 2. Create ClusterIP Services for your applications
- 3. Define Ingress resources with routing rules
- 4. The Ingress Controller configures itself based on your rules
- 5. External DNS points to the Ingress Controller's IP

Common Ingress Controllers

- NGINX Ingress: Most popular, feature-rich
- Traefik: Dynamic configuration, good for microservices
- AWS ALB Ingress: Native integration with AWS ALB
- HAProxy: High performance
- Istio Gateway: Part of service mesh

Ingress provides a powerful, declarative way to manage external access while keeping your services abstracted from networking details.

Kubernetes ConfigMaps and Secrets: Managing Configuration Data

ConfigMaps and Secrets are Kubernetes objects used to decouple configuration data from container images, making applications more portable and secure.

ConfigMaps

Purpose: Store non-sensitive configuration data in key-value pairs.

Use Cases for ConfigMaps:

- 1. Environment variables for applications
- 2. Configuration files (e.g., nginx.conf, application.properties)
- 3. Command-line arguments
- 4. Any non-sensitive app configuration

Example ConfigMap YAML:

```
apiVersion: v1
kind: ConfigMap
metadata:
 name: app-config
data:
 # Key-value pairs
 APP_COLOR: blue
 APP_MODE: prod
 # File-like configuration
  nginx.conf: |
   server {
     listen 80;
     server_name localhost;
      location / {
       proxy_pass http://webapp-service;
   }
```

Using ConfigMap in a Pod:

```
apiVersion: v1
kind: Pod
metadata:
 name: webapp-pod
spec:
 containers:
 - name: webapp
   image: nginx
   envFrom:
   - configMapRef:
       name: app-config # Import all key-value pairs as env vars
   volumeMounts:
   - name: config-volume
     mountPath: /etc/nginx
 volumes:
 - name: config-volume
   configMap:
     name: app-config # Mount specific config files
```

Secrets

PROF P

Purpose: Store sensitive data like passwords, API keys, and TLS certificates in an encrypted form (base64 encoded, not encrypted by default - consider external secret management for production).

Use Cases for Secrets:

- 1. Database credentials
- 2. API tokens/keys
- 3. TLS certificates
- 4. SSH keys
- 5. Any sensitive configuration

Example Secret YAML:

```
apiVersion: v1
kind: Secret
metadata:
   name: db-secret
type: Opaque # Default type for arbitrary data
data:
```

```
# Base64 encoded values (echo -n "value" | base64)

DB_USER: cm9vdA==

DB_PASSWORD: cGFzc3dvcmQxMjM=
```

Using Secret in a Pod:

```
apiVersion: v1
kind: Pod
metadata:
 name: db-pod
spec:
 containers:
  - name: database
   image: mysql
   env:
    - name: MYSQL_ROOT_PASSWORD # Environment variable name
     valueFrom:
        secretKeyRef:
         name: db-secret # Secret name
key: DB_PASSWORD # Key in secret
   volumeMounts:
    - name: cert-volume
      mountPath: /etc/mysql/certs
  volumes:
  - name: cert-volume
    secret:
      secretName: tls-secret
                                 # Mount secret as files
```

Key Differences

Feature	ConfigMap	Secret
Data type	Plain text	Base64 encoded
Security	Non-sensitive data	Sensitive data
Storage	etcd (unencrypted)	etcd (base64, not encrypted)
Use case	Configuration files	Credentials, certificates
Size limitation	1MB per ConfigMap	1MB per Secret

Best Practices

PROF

1. For Secrets:

- Use third-party secret managers (AWS Secrets Manager, HashiCorp Vault) for production
- Enable etcd encryption at rest
- Use Kubernetes RBAC to restrict access
- Consider using SealedSecrets for GitOps workflows

2. For ConfigMaps:

- Use immutable ConfigMaps (immutable: true) for better performance
- Group related configurations together
- Use annotations for metadata/versioning

3. For Both:

- Update references when configs change (may require pod restart)
- Use rolling updates when configuration changes
- Consider using ConfigMap/Secret generators with Kustomize

These mechanisms enable proper separation of configuration from application code, making your Kubernetes applications more secure and portable across environments.

Helm: Kubernetes Package Manager

Helm is the package manager for Kubernetes that simplifies application deployment by:

- Packaging all Kubernetes resources into a single deployable unit (chart)
- Enabling versioning and sharing of applications
- Supporting templating to customize deployments
- Managing dependencies between components

Key Concepts

- 1. **Chart**: A packaged Helm application (collection of YAML templates + metadata)
- 2. **Release**: A deployed instance of a chart
- 3. **Repository**: A collection of shareable charts
- 4. Values: Customizable parameters for the chart

Creating a Helm Chart

1. Initialize a new chart

```
helm create myapp-chart
```

This creates a directory structure:

```
myapp-chart/

— Chart.yaml # Chart metadata
— values.yaml # Default configuration values
— charts/ # Subcharts/dependencies

— templates/ # Kubernetes resource templates
— deployment.yaml
— service.yaml
— ingress.yaml
— helpers.tpl # Template helpers
```

2. Customize the Chart

Edit Chart.yaml:

PROF

```
apiVersion: v2
name: myapp-chart
description: A Helm chart for my application
version: 0.1.0
appVersion: 1.0.0
```

Modify templates/deployment.yaml (example snippet):

```
apiVersion: apps/v1
kind: Deployment
metadata:
   name: {{ include "myapp-chart.fullname" . }}
spec:
   replicas: {{ .Values.replicaCount }}
   template:
     spec:
```

```
containers:
- name: {{ .Chart.Name }}
  image: "{{ .Values.image.repository }}:{{ .Values.image.tag }}"
  ports:
- containerPort: {{ .Values.service.port }}
```

Update values.yaml:

```
replicaCount: 3
image:
    repository: nginx
    tag: stable
    pullPolicy: IfNotPresent
service:
    type: ClusterIP
    port: 80
```

3. Package Dependencies (if any)

```
helm dependency update myapp-chart
```

Deploying the Helm Chart

1. Install the Chart

```
helm install myapp-release ./myapp-chart
```

2. Upgrade with Custom Values

Create a custom values file custom-values.yaml:

```
replicaCount: 5
image:
tag: latest
```

Then upgrade:

```
helm upgrade myapp-release ./myapp-chart -f custom-values.yaml
```

3. Verify Deployment

```
helm list
helm status myapp-release
kubectl get pods
```

4. Rollback (if needed)

```
helm rollback myapp-release 0 # Revert to revision 0
```

Why Helm Simplifies Kubernetes Deployments

1. Templating: Avoids YAML duplication with Go templating

```
env:
{{- range $key, $value := .Values.envVars }}
    - name: {{ $key }}
    value: {{ $value | quote }}
{{- end }}
```

2. Value Management: Single place to control all parameters

```
# values.yaml
resources:
  limits:
    cpu: 500m
    memory: 512Mi
```

3. Release Management: Track versions and rollbacks

```
helm history myapp-release
```

4. Hooks: Execute jobs at specific points in release lifecycle

```
annotations:
"helm.sh/hook": pre-install
```

5. Sharing: Publish to chart repositories

```
helm package myapp-chart
helm push myapp-chart-0.1.0.tgz my-repo
```

Advanced Features

- Library Charts: Reusable chart components
- Chart Tests: Validation for installed charts
- Subcharts: Modular application components
- Post-renderers: Customize manifests after templating

Helm transforms Kubernetes deployments from manual YAML management to a structured, repeatable process with version control and easy customization.

Helm Templates, Values, and Releases Explained with Best Practices

Helm Core Components

- 1. Templates
 - Purpose: Go-templated Kubernetes manifests that become actual resources when deployed
 - Location: templates/ directory in a chart
 - · Features:
 - Use Go template language with Sprig functions
 - Access values from values.yaml via {{ .Values.param }}
 - Include helper functions from <u>helpers.tpl</u>

Example Deployment Template (templates/deployment.yaml):

```
apiVersion: apps/v1
kind: Deployment
metadata:
   name: {{ include "mychart.fullname" . }}
spec:
   replicas: {{ .Values.replicaCount }}
template:
   spec:
    containers:
        - name: {{ .Chart.Name }}
        image: "{{ .Values.image.repository }}:{{ .Values.image.tag }}"
        {{- with .Values.resources }}
        resources:
            {{- toYaml . | nindent 10 }}
        {{- end }}
```

2. Values

- Purpose: Configuration parameters that customize the chart
- **Sources** (in order of precedence):
 - 1. Command-line (--set flags)
 - 2. Custom values files (-f myvalues.yaml)
 - 3. Default values . yaml in chart
 - 4. Parent chart's values. yaml (for subcharts)

Example values.yaml:

```
replicaCount: 3
image:
    repository: nginx
    tag: stable
    pullPolicy: IfNotPresent
resources:
    limits:
        cpu: 200m
        memory: 256Mi
```

3. Releases

- **Definition**: An instance of a chart running in a Kubernetes cluster
- Management:
 - Each helm install creates a new release
 - Releases are versioned (revisions)
 - Stored as Secrets in the cluster by default

Release Lifecycle:

```
helm install \rightarrow helm upgrade \rightarrow helm rollback \rightarrow helm uninstall
```

Helm Chart Best Practices

1. Structure and Organization

```
mychart/

— Chart.yaml  # Required: Chart metadata

— values.yaml  # Required: Default configuration

— charts/  # Optional: Subcharts/dependencies

— templates/  # Required: Kubernetes manifests
```

```
deployment.yaml
   service.yaml
    _helpers.tpl # Template helpers
   - tests/ # Test pods
                 # Documentation
README.md
```

2. Templating Practices

• Use named templates (_helpers.tpl) for reusable components:

```
{{- define "mychart.labels" -}}
app.kubernetes.io/name: {{    .Chart.Name }}
app.kubernetes.io/instance: {{    .Release.Name }}
{{- end }}
```

• Add input validation:

```
{{- if not .Values.image.tag }}
{{- fail "image.tag is required" }}
{{- end }}
```

• Use include over template (preserves scope):

```
labels:
  {{- include "mychart.labels" . | nindent 4 }}
```

3. Values Management

• Document all values in values . yaml with comments:

```
# replicaCount: Number of application instances
replicaCount: 1
```

• Group related values hierarchically:

```
autoscaling:
 enabled: true
 minReplicas: 2
 maxReplicas: 5
```

• Provide sensible defaults but allow overrides:

```
service:
 type: ClusterIP # LoadBalancer for production
 port: 80
```

S

- Never put secrets in values.yaml Use Kubernetes Secrets instead
- Use helm lint to validate charts
- Make charts immutable when possible:

4.	Secur	ity I	Pract	ice
----	-------	-------	-------	-----

```
apiVersion: v2
name: mychart
description: My immutable chart
type: application
```

- 5. Release Management
 - **Use semantic versioning** for charts:

```
version: 1.2.3 # MAJOR.MINOR.PATCH
```

• Implement upgrade hooks for migrations:

```
annotations:
 "helm.sh/hook": pre-upgrade
```

• Test with --dry-run before actual deployment:

```
helm install myapp ./mychart --dry-run --debug
```

- 6. Advanced Techniques
 - Conditional templates:

```
{{- if .Values.ingress.enabled }}
apiVersion: networking.k8s.io/v1
kind: Ingress
{{- end }}
```

• Range over values:

```
{{- range $key, $value := .Values.envVars }}
  - name: {{ $key }}
   value: {{ $value | quote }}
{{- end }}
```

• Template partials for complex objects:

```
{{- define "mychart.resources" }}
{{- if .Values.resources }}
resources:
  {{- toYaml .Values.resources | nindent 4 }}
{{- end }}
{{- end }}
```

Practical Example Workflow

```
1. Create a new chart:
```

```
helm create myapp && rm -rf myapp/templates/*
```

- 2. Add templates (deployment.yaml, service.yaml etc.)
- 3. Define values in values.yaml
- 4. Package and deploy:

```
helm package myapp
helm install myapp ./myapp-0.1.0.tgz --values prod-values.yaml
```

5. Upgrade:

```
helm upgrade myapp ./myapp --set replicaCount=5
```

Helm's templating system combined with these best practices enables you to create maintainable, reusable, and production-grade Kubernetes deployments that can be customized for different environments while keeping your manifests DRY (Don't Repeat Yourself).

Setting Up a CI/CD Pipeline in Jenkins

Here's a complete guide to creating a basic Jenkins pipeline that compiles code and deploys it to a target environment.

Prerequisites

- Jenkins installed (with Pipeline plugin)
- Git repository with your code
- Build tools (Maven, Gradle, npm, etc.) installed on Jenkins agents
- Deployment target access (Kubernetes, server, etc.)

Step 1: Create a New Pipeline Job

- 1. In Jenkins dashboard, click **New Item**
- 2. Enter a name (e.g., "MyApp-CI-CD") and select **Pipeline**
- 3. Click **OK**

Step 2: Configure the Pipeline

PROF O

Option A: Pipeline Script (Jenkinsfile in SCM)

- 1. In your project root, create a Jenkinsfile with the pipeline definition
- 2. In Jenkins job configuration:
 - Select Pipeline script from SCM
 - Choose your SCM (Git, GitHub, etc.)
 - Specify repository URL and credentials
 - Set branch (e.g., main)
 - Specify script path (Jenkinsfile)

Option B: Direct Pipeline Script

- 1. In Jenkins job configuration:
 - Select Pipeline script
 - Paste the pipeline definition directly

Basic Pipeline Example (Java/Maven to Kubernetes)

Here's a complete **Jenkinsfile** example:

```
pipeline {
   agent any
    environment {
        // Define environment variables
       APP_NAME = "my-java-app"
       VERSION = "1.0.${BUILD_NUMBER}"
       DOCKER_IMAGE = "myregistry/${APP_NAME}:${VERSION}"
       KUBE_CONFIG = credentials('kubeconfig') // Stored in Jenkins credentials
   }
    stages {
        stage('Checkout') {
            steps {
                git branch: 'main',
                url: 'https://github.com/your-repo/your-java-app.git'
            }
        }
        stage('Build') {
                sh 'mvn clean package'
                archiveArtifacts artifacts: 'target/*.jar', fingerprint: true
        }
        stage('Test') {
            steps {
                sh 'mvn test'
                junit 'target/surefire-reports/*.xml'
            }
        }
        stage('Build Docker Image') {
            steps {
                script {
                    docker.build(DOCKER_IMAGE)
            }
        }
        stage('Push Image') {
            steps {
                script {
                    docker.withRegistry('https://registry.example.com', 'dockerhub-creds') {
                        docker.image(DOCKER_IMAGE).push()
                }
            }
       }
        stage('Deploy to Kubernetes') {
            steps {
                    kubectl --kubeconfig=${KUBE_CONFIG} set image deployment/${APP_NAME}
${APP_NAME}=${DOCKER_IMAGE}
    }
    post {
        always {
            cleanWs() // Clean workspace
        success {
            slackSend channel: '#deployments',
```

```
message: "SUCCESS: Job ${env.JOB_NAME} build ${env.BUILD_NUMBER}"
}
failure {
    slackSend channel: '#deployments',
    message: "FAILED: Job ${env.JOB_NAME} build ${env.BUILD_NUMBER}"
}
}
}
```

Key Components Explained

1. Pipeline Structure

- agent: Where the pipeline runs (any available agent)
- environment: Shared variables for all stages
- stages: The main workflow steps
- post: Actions after pipeline completes

2. Core Stages

- 1. Checkout: Gets code from version control
- 2. **Build**: Compiles source code (Maven in this case)
- 3. Test: Runs unit tests and reports results
- 4. Docker Build: Creates container image
- 5. **Push Image**: Stores image in registry
- 6. **Deploy**: Updates Kubernetes deployment

3. Advanced Features

- Credentials: Securely stored in Jenkins (kubeconfig, Docker Hub)
- Artifact Archiving: Saves build outputs
- Test Reporting: JUnit test results visualization
- Notifications: Slack integration

Setting Up Required Plugins

Ensure these Jenkins plugins are installed:

- Pipeline
- Docker Pipeline
- Kubernetes CLI
- Git
- JUnit
- Slack Notification
- Credentials Binding

Best Practices

- 1. Use Jenkinsfile in SCM: Store pipeline definition with your code
- 2. Secure Credentials: Never hardcode secrets use Jenkins credentials store
- 3. Parallelize Stages: Run independent stages in parallel
- 4. Implement Rollbacks: Add stage to revert failed deployments
- 5. Parameterize Builds: Allow runtime inputs:

```
parameters {
    choice(name: 'DEPLOY_ENV', choices: ['dev', 'staging', 'prod'], description: 'Target
environment')
    booleanParam(name: 'RUN_TESTS', defaultValue: true, description: 'Run all tests?')
}
```

6. Use Shared Libraries: Reusable pipeline components across projects

Visualizing the Pipeline

Jenkins provides several visualization tools:

- Blue Ocean: Modern UI for pipeline visualization
- Stage View: Traditional stage progress view
- Build History: Timeline of past executions

This pipeline gives you a complete CI/CD workflow from code commit to production deployment, with testing and artifact management built in. You can extend it with additional stages like security scanning, performance testing, or approval gates for production deployments.

Jenkins + Docker Integration: Automated Containerized CI/CD

Here's a complete practical guide to integrating Jenkins with Docker for automating container builds and deployments, including a working example.

Prerequisites

- Jenkins server with Docker plugins installed
- Docker installed on Jenkins agents (or Jenkins server)
- Docker Hub/registry credentials configured in Jenkins
- Kubernetes cluster (for deployment, optional)

Step 1: Configure Jenkins for Docker

1. Install required plugins:

- Docker
- Docker Pipeline
- Docker Build Step

2. Add Docker credentials in Jenkins:

- \circ Go to Credentials \rightarrow System \rightarrow Global credentials
- Add username/password for your Docker registry

Step 2: Create a Docker-Enabled Pipeline

Here's a complete Jenkinsfile example that:

- 1. Builds a Node.js application
- 2. Creates a Docker image
- 3. Pushes to Docker Hub
- 4. Deploys to Kubernetes

pipeline {

}

```
agent {
    docker {
        image 'node:14-alpine'
        reuseNode true
```

```
args '-u root:root' // Needed for file permissions
    }
}
environment {
   APP_NAME = "nodejs-app"
   DOCKER_IMAGE = "your-dockerhub/${APP_NAME}:${BUILD_NUMBER}"
   KUBE_NAMESPACE = "default"
    // Credentials from Jenkins
    DOCKER_CREDS = credentials('dockerhub-creds')
    KUBECONFIG = credentials('kubeconfig')
```

```
stages {
        stage('Checkout') {
            steps {
                checkout scm
       }
        stage('Install Dependencies') {
            steps {
                sh 'npm install'
       }
        stage('Run Tests') {
            steps {
                sh 'npm test'
       }
        stage('Build Docker Image') {
            agent any // Needs Docker daemon
            steps {
                script {
                    docker.build(DOCKER_IMAGE)
            }
        }
        stage('Push Image') {
            steps {
                script {
                    docker.withRegistry('https://registry.hub.docker.com', DOCKER_CREDS) {
                        docker.image(DOCKER_IMAGE).push()
                }
            }
       }
        stage('Deploy to Kubernetes') {
            steps {
                    kubectl --kubeconfig=${KUBECONFIG} -n ${KUBE_NAMESPACE} \
                    set image deployment/${APP_NAME} ${APP_NAME}=${DOCKER_IMAGE}
                0.00
            }
       }
   }
    post {
        always {
            cleanWs()
        success {
            slackSend(color: 'good', message: "SUCCESS: ${JOB_NAME} #${BUILD_NUMBER}")
        failure {
            slackSend(color: 'danger', message: "FAILED: ${JOB_NAME} #${BUILD_NUMBER}")
       }
   }
}
```

Key Integration Points Explained

1. Docker Build Agent

```
agent {
    docker {
        image 'node:14-alpine' // Uses Node.js container for build
        args '-u root:root'
    }
}
```

- Runs build stages inside a containerized environment
- Ensures consistent build tools across executions

2. Docker Image Building

```
script {
   docker.build(DOCKER_IMAGE)
}
```

- Requires a Dockerfile in your project root
- Example Dockerfile for Node.js app:

```
FROM node:14-alpine
WORKDIR /app
COPY package*.json ./
RUN npm install --production
COPY . .

EXPOSE 3000
CMD ["npm", "start"]
```

3. Secure Registry Push

```
docker.withRegistry('https://registry.hub.docker.com', DOCKER_CREDS) {
   docker.image(DOCKER_IMAGE).push()
}
```

- Uses Jenkins-stored credentials
- Supports private registries (Artifactory, ECR, GCR, etc.)

4. Kubernetes Deployment

```
sh """

kubectl set image deployment/${APP_NAME} ${APP_NAME}=${DOCKER_IMAGE}
"""
```

- Assumes existing Kubernetes deployment
- Alternative: kubectl apply -f k8s-deployment.yaml

Advanced Integration Patterns

1. Multi-Architecture Builds

```
stage('Build Multi-Arch Image') {
   steps {
     sh """
        docker buildx build --platform linux/amd64,linux/arm64 \
        -t ${DOCKER_IMAGE} --push .
        """
```

```
}
```

2. Docker Compose Testing

```
stage('Integration Tests') {
    steps {
        sh 'docker-compose -f docker-compose.test.yml up --build --abort-on-container-exit'
    }
}
```

3. Parameterized Builds

```
parameters {
   choice(name: 'ENVIRONMENT', choices: ['dev', 'staging', 'prod'])
   string(name: 'IMAGE_TAG', defaultValue: 'latest')
}
```

4. Docker Content Trust

```
environment {
   DOCKER_CONTENT_TRUST = '1'
}
```

Best Practices

1. Use Docker-in-Docker (DinD) for isolated builds:

```
agent {
    docker {
        image 'docker:dind'
        args '--privileged -v /var/run/docker.sock:/var/run/docker.sock'
    }
}
```

2. Cache dependencies between builds:

3. Scan images for vulnerabilities:

```
stage('Security Scan') {
   steps {
     sh 'docker scan ${DOCKER_IMAGE}'
   }
}
```

4. Clean up old images:

```
post {
    always {
       sh 'docker system prune -f || true'
    }
}
```

Monitoring the Pipeline

- 1. Blue Ocean View: Visual pipeline progress
- 2. Console Output: Detailed build logs
- 3. Docker Hub: Verify pushed images
- 4. Kubernetes: Check deployment status

```
kubectl get pods -w
kubectl describe deployment/nodejs-app
```

This integration provides a complete workflow from code commit to production deployment, leveraging Docker for consistent builds and environments throughout the pipeline.

Configuring GitHub Webhooks for Jenkins Auto-Triggered Builds

Here's a complete step-by-step guide to setting up GitHub webhooks that automatically trigger Jenkins builds when code is pushed:

Prerequisites

- Jenkins server accessible via HTTPS (GitHub requires HTTPS for webhooks)
- GitHub repository admin access
- Jenkins GitHub Plugin installed (GitHub plugin and GitHub API plugin)

Step 1: Configure Jenkins for GitHub Integration

- 1. Install Required Plugins
 - Manage Jenkins → Manage Plugins → Available
 - Install:
 - GitHub plugin
 - GitHub API plugin
 - (Optional) Pipeline: GitHub Plugin
- 2. Set Up GitHub Server in Jenkins
 - 1. Go to Manage Jenkins → Configure System
 - 2. Find **GitHub** section
 - 3. Add GitHub Server:
 - Name: GitHub
 - APIURL: https://api.github.com
 - 4. Add credentials:
 - Kind: Secret text
 - Secret: [GitHub personal access token] (with repo scope)
 - ID: github-token

Step 2: Configure the Jenkins Job

For Freestyle Projects:

1. Create/Edit job → **Configure**

- 2. Under **Source Code Management**:
 - Select Git
 - Enter Repository URL: https://github.com/your-username/your-repo.git
- 3. Under Build Triggers:
 - Check GitHub hook trigger for GITScm polling

For Pipeline Projects:

1. Use this trigger definition in your Jenkinsfile:

```
triggers {
   githubPush()
}
```

Or for more control:

```
triggers {
    GitHubPushTrigger {
        branchesToBuild = ['main', 'develop']
    }
}
```

Step 3: Set Up GitHub Webhook

- 1. Go to your GitHub repository → **Settings** → **Webhooks** → **Add webhook**
- 2. Configure the webhook:
 - Payload URL: https://your-jenkins-server/github-webhook/
 - Content type: application/json
 - Secret: [Optional] Add a secret that matches Jenkins configuration
 - Which events:
 - For push events: Just the push event
 - For all events: **Send me everything**
- 3. Click Add webhook

Step 4: Verify the Connection

- 1. In GitHub, check the webhook's recent deliveries
- 2. In Jenkins, check:
 - ullet Manage Jenkins ullet System Log for GitHub plugin logs
 - Job's **Build History** for triggered builds

Troubleshooting Tips

Common Issues and Fixes:

- 1. Webhook not triggering:
 - Verify Jenkins URL is accessible from GitHub
 - Check Jenkins logs (/var/log/jenkins/jenkins.log)
 - Test with curl -X POST https://your-jenkins-server/github-webhook/
- 2. 403 Forbidden errors:
 - Ensure GitHub token has proper permissions
 - Verify webhook secret matches (if used)
- 3. Builds not starting:
 - Check "GitHub hook trigger" is enabled in job config
 - Verify SCM polling is disabled (conflicts with webhooks)

Advanced Configuration

1. Filtering Specific Branches

In your Jenkinsfile:

```
triggers {
    githubPush()
}
pipeline {
    agent any
    stages {
        starge('Build') {
            when {
                branch 'main' // Only build on main branch pushes
            }
            steps {
                sh 'make build'
            }
        }
    }
}
```

2. Using Webhook Secrets

1. Generate a secret:

```
openssl rand -hex 20
```

- 2. Add to Jenkins GitHub server config
- 3. Use same secret in GitHub webhook

3. Multibranch Pipelines

- 1. Create **New Item** → **Multibranch Pipeline**
- 2. Configure branch sources to auto-discover branches/PRs
- 3. Webhooks will automatically trigger builds for all branches

Security Best Practices

PROF

- 1. **Use HTTPS**: GitHub requires HTTPS for webhooks
- 2. **Restrict IPs**: Configure GitHub webhook to only allow Jenkins server IP
- 3. Use Secrets: Always use webhook secrets for verification
- 4. Limit Permissions: GitHub token only needs repo scope
- 5. Rotate Secrets: Periodically update webhook and API tokens

This setup creates a fully automated CI pipeline where any push to GitHub triggers the corresponding Jenkins job immediately, enabling true continuous integration.

Jenkins Pipeline as Code: Declarative vs Scripted Pipelines

Jenkins Pipeline as Code is an infrastructure-as-code approach that defines CI/CD pipelines using code (typically Groovy) stored in version control, rather than configuring jobs through Jenkins' UI.

Pipeline as Code Benefits

- Version-controlled pipeline definitions
- Reproducible builds across environments
- Code review for pipeline changes
- Reusable components via shared libraries

• Auditable change history

Declarative Pipeline

The newer, simpler syntax with opinionated structure, ideal for most use cases.

Basic Structure:

```
pipeline {
   agent any
   stages {
        stage('Build') {
            steps {
                // Build steps
        }
        stage('Test') {
           steps {
                // Test steps
        }
    }
    post {
       always {
           // Cleanup steps
   }
}
```

Complete Example:

```
// Jenkinsfile (Declarative)
pipeline {
    agent {
        docker {
           image 'maven:3.8.6-jdk-11'
            args '-v $HOME/.m2:/root/.m2' // Cache Maven artifacts
       }
    }
    options {
        timeout(time: 1, unit: 'HOURS')
        buildDiscarder(logRotator(numToKeepStr: '10'))
   }
    environment {
       VERSION = "1.0.${BUILD_NUMBER}"
       DOCKER_IMAGE = "my-registry/app:${VERSION}"
    }
    stages {
        stage('Checkout') {
            steps {
                checkout scm
        }
        stage('Build') {
            steps {
                sh 'mvn clean package'
                archiveArtifacts artifacts: 'target/*.jar', fingerprint: true
            }
       }
```

```
stage('Test') {
            steps {
                sh 'mvn test'
                junit 'target/surefire-reports/*.xml'
            }
       }
        stage('Docker Build') {
            when {
                branch 'main'
            }
            steps {
                script {
                    docker.build(DOCKER_IMAGE)
        }
   }
    post {
        success {
            slackSend channel: '#builds',
                   message: "SUCCESS: ${JOB_NAME} #${BUILD_NUMBER}"
        }
        failure {
            slackSend channel: '#builds',
                    message: "FAILED: ${JOB_NAME} #${BUILD_NUMBER}"
       }
    }
}
```

Scripted Pipeline

The original flexible Groovy-based syntax for advanced use cases.

Basic Structure:

```
node {
    stage('Build') {
        // Build steps
    }
    stage('Test') {
        // Test steps
    }
}
```

Complete Example:

```
// Jenkinsfile (Scripted)
def dockerImage = null

node('docker-agent') {
    checkout scm

stage('Build') {
        docker.image('maven:3.8.6-jdk-11').inside('-v $HOME/.m2:/root/.m2') {
            sh 'mvn clean package'
        }
        archiveArtifacts artifacts: 'target/*.jar', fingerprint: true
}

stage('Test') {
        docker.image('maven:3.8.6-jdk-11').inside {
```

```
sh 'mvn test'
        junit 'target/surefire-reports/*.xml'
    }
    if (env.BRANCH_NAME == 'main') {
        stage('Docker Build') {
            dockerImage = docker.build("my-registry/app:1.0.${BUILD_NUMBER}")
        stage('Deploy') {
            docker.withRegistry('https://my-registry', 'docker-creds') {
                dockerImage.push()
            sh "kubectl set image deployment/app app=my-registry/app:1.0.${BUILD_NUMBER}"
        }
    }
}
// Post-build actions
if (currentBuild.result == 'SUCCESS') {
    slackSend channel: '#builds',
             message: "SUCCESS: ${JOB_NAME} #${BUILD_NUMBER}"
} else {
    slackSend channel: '#builds',
            message: "FAILED: ${JOB_NAME} #${BUILD_NUMBER}"
}
```

Key Differences

Feature	Declarative	Scripted
Syntax	Structured with predefined sections	Flexible Groovy code
Learning Curve	Easier	Steeper (requires Groovy knowledge)
Validation	Early syntax validation	Runtime error checking
When to Use	Most standard CI/CD pipelines	Complex logic/advanced workflows
Directives	Uses pipeline, stages, etc.	Uses Groovy constructs
Error Handling	Built-in post conditions	Manual try/catch blocks
Restartability	Built-in checkpoint support	Manual implementation

Best Practices for Both Styles

- 1. Store in SCM: Always keep Jenkinsfile in version control
- 2. Start Simple: Begin with declarative, switch to scripted only when needed
- 3. Use Shared Libraries: For reusable functions across pipelines
- 4. Parameterize Builds: Allow runtime configuration

```
parameters {
    string(name: 'DEPLOY_ENV', defaultValue: 'staging')
}
```

5. Secure Credentials: Use Jenkins credential store

```
environment {
    AWS_ACCESS_KEY_ID = credentials('aws-access-key')
}
```

6. Clean Up: Always clean workspace in post-build

```
post {
    always {
      cleanWs()
    }
}
```

When to Choose Which

Use Declarative When:

- You need a simple, readable pipeline
- You want built-in syntax checking
- Your workflow fits standard CI/CD patterns
- You're new to Jenkins Pipelines

Use Scripted When:

- You need complex logic (loops, conditionals)
- You require advanced error handling
- You're integrating with external systems
- · You need dynamic pipeline generation

Both approaches can coexist - you can call scripted code from declarative pipelines using the script block when needed.

Git Branching Strategies and CI/CD Integration: GitFlow in Practice

GitFlow is a branching model that provides a structured approach to managing features, releases, and hotfixes in a way that enhances CI/CD processes. Here's how it works with CI/CD pipelines and a practical setup example.

How GitFlow Enhances CI/CD

1. Clear Environment Mapping:

- develop → Staging environment
- main → Production environment
- feature/* → Ephemeral test environments

2. Automated Quality Gates:

- Feature branches run unit tests
- Develop branch runs integration tests
- Release branches run full regression tests

3. Controlled Releases:

- Release branches enable gradual rollout
- Hotfix branches allow emergency patches

4. Parallel Development:

- Multiple features can progress simultaneously
- Release preparation doesn't block new development

GitFlow CI/CD Setup Example

1. Branch Structure

```
main - Production code (always deployable)
release/* - Release preparation branches
```

```
develop - Integration branch (nightly builds)
feature/* - Feature development branches
hotfix/* - Critical bug fixes
```

2. Jenkins Pipeline Setup (Declarative)

```
// Jenkinsfile at repository root
pipeline {
   agent any
   options {
       skipDefaultCheckout true
   }
    stages {
        stage('Checkout & Detect Branch Type') {
            steps {
                checkout scm
                script {
                    // Determine branch type
                    if (env.BRANCH_NAME == 'main') {
                        env.BRANCH_TYPE = 'production'
                    } else if (env.BRANCH_NAME == 'develop') {
                        env.BRANCH_TYPE = 'staging'
                    } else if (env.BRANCH_NAME.startsWith('release/')) {
                        env.BRANCH_TYPE = 'release'
                    } else if (env.BRANCH_NAME.startsWith('feature/')) {
                        env.BRANCH_TYPE = 'feature'
                    } else if (env.BRANCH_NAME.startsWith('hotfix/')) {
                        env.BRANCH_TYPE = 'hotfix'
                    }
                }
            }
       }
        stage('Build') {
            steps {
                sh 'mvn clean package'
                stash name: 'artifacts', includes: 'target/*.jar'
            }
        }
        stage('Unit Tests') {
            steps {
                sh 'mvn test'
                junit 'target/surefire-reports/*.xml'
            }
        }
        stage('Integration Tests') {
            when {
                expression {
                    env.BRANCH_TYPE == 'develop' ||
                    env.BRANCH_TYPE == 'release'
                }
            }
            steps {
                sh 'mvn verify -Pintegration-tests'
                archiveArtifacts artifacts: 'target/*.war', fingerprint: true
            }
        stage('Deploy') {
            steps {
                script {
                    switch(env.BRANCH_TYPE) {
                        case 'feature':
```

```
// Deploy to feature env
                          sh "kubectl apply -f
break
                      case 'develop':
                          // Deploy to staging
                          sh 'kubectl apply -f k8s/overlays/staging'
                          break
                      case 'release':
                      case 'hotfix':
                          // Deploy to pre-prod
                          sh 'kubectl apply -f k8s/overlays/pre-prod'
                          break
                      case 'production':
                          // Deploy to production
                          sh 'kubectl apply -f k8s/overlays/production'
                  }
              }
           }
       }
       stage('Approval') {
           when {
               anyOf {
                  branch 'release/*'
                  branch 'hotfix/*'
               }
           }
           steps {
               timeout(time: 1, unit: 'HOURS') {
                  input message: "Deploy to production?", ok: "Confirm"
           }
       }
   }
   post {
       success {
           script {
               if (env.BRANCH_TYPE == 'feature') {
                  slackSend channel: '#features',
                           message: "Feature ${env.BRANCH_NAME} ready for review"
              }
           }
       }
       always {
           cleanWs()
       }
   }
}
```

3. GitFlow Workflow with CI/CD

Feature Development:

PROF

- 1. Developer creates feature branch: git checkout -b feature/new-payment develop
- 2. CI pipeline runs on every push:
 - Builds code
 - Runs unit tests
 - Deploys to ephemeral environment (feature.new-payment.example.com)

Release Preparation:

- 1. Create release branch: git checkout -b release/1.2 develop
- 2. CI pipeline:

- Runs full integration tests
- Deploys to pre-production
- Requires manual approval for production

Hotfix Workflow:

- 1. Create from main: git checkout -b hotfix/urgent-fix main
- 2. CI pipeline:
 - Runs critical path tests only
 - Fast-tracks to production after approval

4. Kubernetes Overlay Structure

```
k8s/

├── base/  # Common configurations

├── overlays/

├── feature/  # Feature-specific configs

├── feature-new-payment/

├── staging/  # Develop branch configs

├── pre-prod/  # Release branch configs

├── production/  # Main branch configs
```

5. Automated Version Bumping

Benefits of This Setup

- 1. **Environment Consistency**: Each branch type maps to a specific environment
- 2. **Quality Enforcement**: Test requirements escalate toward production
- 3. Traceability: Every production deploy comes from an explicit release branch
- 4. Parallel Workflows: Features and releases can progress simultaneously
- 5. **Emergency Path**: Hotfixes bypass normal workflow when needed

Advanced Enhancements

1. Automated PR Gates:

```
stage('PR Validation') {
   when {
      changeRequest()
   }
   steps {
      sh 'mvn verify -Pquick-tests'
   }
}
```

2. Feature Environment Teardown:

```
post {
    cleanup {
        script {
            if (env.BRANCH_TYPE == 'feature') {
                sh "kubectl delete -f
            k8s/overlays/feature/${env.BRANCH_NAME.replaceAll('/', '-')}"
            }
        }
    }
}
```

3. Release Notes Generation:

```
stage('Generate Release Notes') {
    when {
        branch 'release/*'
    }
    steps {
        sh """
            git log --pretty=format:"- %s" develop..HEAD > release-notes.md
        """
            archiveArtifacts artifacts: 'release-notes.md'
    }
}
```

This GitFlow+CI/CD integration provides a robust framework that balances development velocity with production stability, ensuring code progresses through well-defined quality gates before reaching users.

Setting Up Prometheus for Kubernetes Monitoring

Prometheus is the de facto standard for monitoring Kubernetes applications. Here's a comprehensive guide to setting it up and configuring key metrics and alerts.

Installation Methods

1. Using Helm (Recommended)

PROF

```
# Add Prometheus community charts
helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
helm repo update

# Install kube-prometheus stack (includes Grafana)
helm install prometheus prometheus-community/kube-prometheus-stack \
    --namespace monitoring \
    --create-namespace \
    --set prometheus.prometheusSpec.serviceMonitorSelectorNilUsesHelmValues=false
```

2. Using Prometheus Operator

```
kubectl apply -f https://raw.githubusercontent.com/prometheus-operator/kube-
prometheus/main/manifests/setup.yaml
kubectl apply -f https://raw.githubusercontent.com/prometheus-operator/kube-
prometheus/main/manifests/
```

Key Components

- 1. **Prometheus Server**: Time-series database and evaluation engine
- 2. **Alertmanager**: Handles alerts routing and deduplication
- 3. **Grafana**: Visualization dashboard (included in kube-prometheus-stack)
- 4. Exporters:
 - kube-state-metrics: Kubernetes object state metrics
 - node-exporter: Node-level hardware metrics
 - application-specific exporters

Key Kubernetes Metrics to Monitor

Cluster-Level Metrics

Metric	Description	Alert Threshold Example	
kube_node_status_condition	Node health conditions	condition="Ready", status!="true"	
kube_pod_status_phase	Pod status phases	phase!="Running" for 5m	
kube_deployment_status_replicas_unavailable	Unavailable replicas	> 0 for 10m	
kube_node_status_allocatable_cpu_cores	Available CPU	<pre>(sum by (node) - sum by (node) (rate(container_cpu_usage_seconds_total[5m])) / sum by (node) (kube_node_status_allocatable_cpu_cores) > 0.9</pre>	

Node-Level Metrics

Metric	Description	Alert Threshold
node_memory_MemAvailable_bytes	Available memory	< 10% total memory
node_cpu_seconds_total	CPU usage	> 90% utilization
node_filesystem_avail_bytes	Disk space	< 15% free space
node_network_receive_bytes_total	Network traffic	Spikes detection

Application-Level Metrics

PROF

Metric	Description	Alert Threshold
http_requests_total	Request count	Sudden drops
http_request_duration_seconds	Latency	99th percentile > 1s
container_memory_working_set_bytes	Memory usage	> 90% of limit
container_cpu_usage_seconds_total	CPU usage	> 80% of limit

Configuring Service Monitoring

1. ServiceMonitor for Your Application

```
apiVersion: monitoring.coreos.com/v1
kind: ServiceMonitor
metadata:
   name: webapp-monitor
   namespace: your-app-ns
spec:
   selector:
   matchLabels:
```

```
app: webapp
endpoints:
- port: web
path: /metrics
interval: 30s
```

2. PodMonitor for Stateful Applications

```
apiVersion: monitoring.coreos.com/v1
kind: PodMonitor
metadata:
   name: redis-monitor
spec:
   selector:
    matchLabels:
        app: redis
podMetricsEndpoints:
        - port: metrics
```

Alerting Rules Configuration

1. Custom Alert Rules

```
apiVersion: monitoring.coreos.com/v1
kind: PrometheusRule
metadata:
 name: custom-rules
 namespace: monitoring
 groups:
  - name: node-alerts
   rules:
    - alert: HighNodeCPU
     expr: 100 - (avg by (instance) (irate(node_cpu_seconds_total{mode="idle"}[5m])) * 100 > 90
     for: 10m
     labels:
       severity: warning
      annotations:
       summary: "High CPU usage on {{ $labels.instance }}"
       description: "CPU usage is {{ $value }}%"
  - name: application-alerts
   rules:
    - alert: HighLatency
     expr: histogram_quantile(0.99, sum(rate(http_request_duration_seconds_bucket[5m])) by (le,
service) > 1
     for: 5m
      labels:
       severity: critical
```

Key Alert Categories

1. Resource Alerts:

- CPU/Memory/Disk pressure
- Node not ready
- Persistent volume claims full

2. Application Alerts:

```
High error rates (rate(http_requests_total{status_code=~"5.."}[5m]))
```

- Increased latency
- Dead letter queues

3. Business Metrics:

- Transaction failures
- Order processing delays
- API rate limiting

Accessing the Dashboards

1. Prometheus UI:

```
kubectl port-forward -n monitoring svc/prometheus-operated 9090
```

Access at http://localhost:9090

2. Grafana:

```
kubectl port-forward -n monitoring svc/prometheus-grafana 3000:80
```

Default credentials: admin/prom-operator

Advanced Configuration

1. Long-Term Storage with Thanos

```
# values.yaml for kube-prometheus-stack
prometheus:
  prometheusSpec:
    thanos:
    objectStorageConfig:
        key: thanos.yaml
        name: thanos-objstore-config
```

2. Custom Scraping Config

```
additionalScrapeConfigs:
    job_name: 'external-service'
    static_configs:
    targets: ['external.service:9115']
```

3. Recording Rules

```
groups:
    name: recording-rules
    rules:
    - record: cluster:node_cpu:avg_rate5m
        expr: avg by (cluster) (rate(node_cpu_seconds_total[5m]))
```

Best Practices

- 1. Labeling Strategy: Consistent labels across services
- 2. Metric Cardinality: Avoid high-cardinality labels
- 3. Alert Fatigue: Set meaningful thresholds and durations

4. **SLO Monitoring**: Track error budgets

```
(sum(rate(http_requests_total{status_code!~"5.."}[7d])) /
sum(rate(http_requests_total[7d])))
```

5. Multi-Cluster: Use Thanos or Cortex for federation

This setup provides comprehensive visibility into your Kubernetes applications with meaningful alerts to proactively address issues before they impact users.

Setting Up EFK Stack for Kubernetes Logging

The EFK (Elasticsearch, Fluentd, Kibana) stack is a powerful solution for centralized logging in Kubernetes environments. Here's a complete guide to deploying and configuring it:

Prerequisites

- Kubernetes cluster (v1.19+ recommended)
- kubectl configured with cluster access
- Helm v3+ installed
- Minimum 4GB RAM available in cluster

Deployment Methods

Method 1: Using Helm (Recommended)

1. Add Helm Repositories:

```
helm repo add elastic https://helm.elastic.co
helm repo add fluent https://fluent.github.io/helm-charts
helm repo update
```

2. Install Elasticsearch:

```
helm install elasticsearch elastic/elasticsearch \
--namespace logging \
--create-namespace \
--set replicas=3 \
--set minimumMasterNodes=2 \
--set resources.requests.memory=4Gi \
--set resources.limits.memory=8Gi
```

3. Install Fluentd:

```
helm install fluentd fluent/fluentd \
--namespace logging \
--set elasticsearch.host=elasticsearch-master.logging.svc.cluster.local \
--set elasticsearch.port=9200 \
--set resources.requests.memory=512Mi \
--set resources.limits.memory=1Gi
```

4. Install Kibana:

```
helm install kibana elastic/kibana \
--namespace logging \
```

```
--set service.type=LoadBalancer \
--set elasticsearchHosts=http://elasticsearch-master.logging.svc.cluster.local:9200
```

Method 2: Manual YAML Deployment

1. Create Namespace:

```
kubectl create namespace logging
```

2. Deploy Elasticsearch StatefulSet (elasticsearch.yaml):

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
 name: elasticsearch
 namespace: logging
spec:
  serviceName: elasticsearch
  replicas: 3
  selector:
   matchLabels:
     app: elasticsearch
  template:
    metadata:
      labels:
       app: elasticsearch
    spec:
      containers:
      - name: elasticsearch
        image: docker.elastic.co/elasticsearch/elasticsearch:7.17.3
        ports:
        - containerPort: 9200
         name: http
        - containerPort: 9300
          name: transport
        volumeMounts:
        - name: data
          mountPath: /usr/share/elasticsearch/data
        - name: discovery.type
         value: single-node # For production, use "zen" with proper config
        - name: ES_JAVA_OPTS
         value: "-Xms2g -Xmx2g"
      volumes:
      - name: data
        emptyDir: {}
```

3. Deploy Fluentd DaemonSet (fluentd.yaml):

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
    name: fluentd
    namespace: logging
spec:
    selector:
    matchLabels:
        app: fluentd
template:
    metadata:
    labels:
```

```
app: fluentd
spec:
 containers:
  - name: fluentd
   image: fluent/fluentd-kubernetes-daemonset:v1.14.6-debian-elasticsearch7-1.0
   - name: FLUENT_ELASTICSEARCH_HOST
     value: "elasticsearch.logging.svc.cluster.local"
    - name: FLUENT_ELASTICSEARCH_PORT
     value: "9200"
   volumeMounts:
    - name: varlog
     mountPath: /var/log
    - name: varlibdockercontainers
     mountPath: /var/lib/docker/containers
     readOnly: true
 volumes:
  - name: varlog
   hostPath:
     path: /var/log
  - name: varlibdockercontainers
   hostPath:
     path: /var/lib/docker/containers
```

4. Deploy Kibana Deployment (kibana.yaml):

```
apiVersion: apps/v1
kind: Deployment
metadata:
 name: kibana
  namespace: logging
spec:
  replicas: 1
  selector:
   matchLabels:
     app: kibana
  template:
    metadata:
      labels:
        app: kibana
    spec:
     containers:
      - name: kibana
        image: docker.elastic.co/kibana/kibana:7.17.3
        ports:
        - containerPort: 5601
        env:
        - name: ELASTICSEARCH_HOSTS
          value: "http://elasticsearch.logging.svc.cluster.local:9200"
```

Configuration Details

Fluentd Configuration (Customize as needed)

Create a ConfigMap for custom Fluentd configuration (fluentd-config.yaml):

```
apiVersion: v1
kind: ConfigMap
metadata:
   name: fluentd-config
   namespace: logging
data:
   fluent.conf: |
```

```
<source>
  @type tail
  path /var/log/containers/*.log
  pos_file /var/log/fluentd-containers.log.pos
  tag kubernetes.*
  read_from_head true
  <parse>
    @type json
    time_format %Y-%m-%dT%H:%M:%S.%NZ
  </parse>
</source>
<filter kubernetes.**>
 @type kubernetes_metadata
</filter>
<match kubernetes.**>
 @type elasticsearch
 host elasticsearch.logging.svc.cluster.local
  port 9200
  logstash_format true
  logstash_prefix fluentd
  include_tag_key true
  type_name _doc
</match>
```

Elasticsearch Resource Requirements

For production environments, consider these settings:

```
resources:
    requests:
    cpu: "1000m"
    memory: "4Gi"
limits:
    cpu: "2000m"
    memory: "8Gi"
```

Verification and Access

1. Check Pod Status:

```
kubectl get pods -n logging
```

2. Access Kibana:

```
kubectl port-forward svc/kibana-kibana 5601:5601 -n logging
```

Open http://localhost:5601

3. Verify Logs in Elasticsearch:

```
kubectl exec -it elasticsearch-0 -n logging -- curl http://localhost:9200/_cat/indices?v
```

Production Considerations

1. Persistent Storage:

```
volumeClaimTemplates:
- metadata:
    name: data
spec:
    accessModes: [ "ReadWriteOnce" ]
    storageClassName: "standard"
    resources:
    requests:
    storage: 50Gi
```

2. Security:

- Enable Elasticsearch security features (x-pack)
- Use network policies to restrict access
- Implement TLS between components

3. Log Retention:

Configure Index Lifecycle Management (ILM) policies in Elasticsearch

4. Multi-Line Logs:

Add Fluentd multiline parsing for stack traces:

```
<filter kubernetes.**>
  @type concat
  key log
  multiline_start_regexp /^\d{4}-\d{2}-\d{2}/
</filter>
```

5. Custom Log Processing:

```
<filter app.**>
@type record_transformer
<record>
    service_name ${record.dig("kubernetes", "labels", "app")}
</record>
</filter>
```

Troubleshooting

No Logs in Kibana:

- Check Fluentd logs: kubectl logs fluentd-xxxx -n logging
- Verify Elasticsearch connection from Fluentd
- Check index patterns in Kibana

2. Performance Issues:

- Increase Fluentd buffer size
- Scale Elasticsearch nodes
- Optimize Elasticsearch shard count

This EFK stack setup provides a robust logging solution that collects, processes, stores, and visualizes logs from all Kubernetes pods and nodes. The Helm method is recommended for most deployments as it simplifies maintenance and upgrades.

Kubernetes Load Balancing: Ingress vs. Service Mesh vs. External Load Balancers

Kubernetes offers multiple approaches to load balancing traffic, each serving different needs in the application architecture. Here's a comprehensive comparison of the three main methods:

1. Kubernetes Ingress

Purpose: HTTP/HTTPS layer 7 routing and load balancing

Implementation:

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
 name: app-ingress
 annotations:
    nginx.ingress.kubernetes.io/load-balance: "ewma"
 rules:
  - host: myapp.example.com
   http:
     paths:
      - path: /
       pathType: Prefix
       backend:
         service:
           name: app-service
            port:
             number: 80
```

Characteristics:

- Pros:
 - Path-based and host-based routing
 - TLS termination
 - Cost-efficient (single LB IP)
 - Annotation-based customization
- Cons:
 - HTTP/HTTPS only (L7)
 - Limited TCP/UDP support
 - No fine-grained traffic control

Use Cases:

PROF

- Web applications needing URL routing
- API gateways
- TLS termination for multiple services

2. Service Mesh (Istio, Linkerd)

Purpose: Fine-grained L7 traffic management with observability

Istio Example:

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
   name: app-virtual-service
spec:
   hosts:
   - "myapp.example.com"
http:
   - route:
    - destination:
```

```
host: app-service
subset: v1
weight: 90
- destination:
host: app-service
subset: v2
weight: 10
```

Characteristics:

- Pros:
 - Advanced traffic splitting (canary, blue-green)
 - Circuit breaking
 - Automatic mTLS
 - Rich observability (metrics, tracing)
 - Service-to-service load balancing
- Cons:
 - Complex setup
 - Performance overhead (sidecar proxies)
 - Steeper learning curve

Use Cases:

- Microservices architectures
- Progressive rollouts
- Zero-trust networks
- Complex routing scenarios

3. External Load Balancers (Cloud Providers)

Purpose: Direct L3/L4 traffic distribution

Implementation:

```
apiVersion: v1
kind: Service
metadata:
   name: app-loadbalancer
spec:
   type: LoadBalancer
ports:
   - port: 443
     targetPort: 8443
selector:
   app: my-app
```

• Pros:

Characteristics:

- Simple configuration
- Works with any protocol (TCP/UDP)
- Cloud-native integrations (AWS ALB, GCP LB)
- High performance
- Cons:
 - Cost (per-LB pricing)
 - Limited to L4 features
 - No application awareness
 - Multiple IPs needed for multiple services

Use Cases:

- Non-HTTP services (gRPC, databases)
- Performance-critical applications
- Legacy applications

Comparative Analysis

Feature	Ingress	Service Mesh	External LB
Layer	L7 (HTTP/HTTPS)	L4-L7	L3-L4
Routing	Host/path-based	Advanced policies	Port-based
TLS	Termination	End-to-end mTLS	Passthrough
Traffic Control	Basic	Fine-grained	None
Observability	bservability Basic metrics		Cloud provider
Performance	Good	Overhead	Excellent
Cost	Low	Moderate	High (per-LB)
Complexity	Moderate	High	Low

Hybrid Approaches

1. Ingress + Service Mesh

```
graph LR
   A[User] --> B(Ingress LB)
   B --> C[Istio Ingress Gateway]
   C --> D[Service Mesh]
   D --> E[Pods]
```

2. Global LB + Ingress

```
graph LR
   A[User] --> B[Cloud CDN]
   B --> C[Regional LBs]
   C --> D[Cluster Ingress]
   D --> E[Services]
```

Best Practices

PROF

- 1. **Start Simple**: Begin with Ingress for web apps
- 2. Adopt Service Mesh Gradually: Only when you need advanced features
- 3. Use External LBs Judiciously: Reserve for specific performance needs
- 4. Layer Solutions:
 - Global LB for geo-distribution
 - Ingress for HTTP routing
 - Service mesh for internal traffic
- 5. Monitor Performance: Watch for latency introduced by each layer

Emerging Patterns

1. Gateway API: Next-gen Ingress (in beta)

```
apiVersion: gateway.networking.k8s.io/v1beta1
kind: HTTPRoute
metadata:
   name: app-route
spec:
```

```
rules:
- matches:
- path:
    type: PathPrefix
    value: /v2
filters:
- type: RequestMirror
    requestMirror:
    backendRef:
    name: app-v2
    port: 80
```

2. eBPF-based Load Balancing: Cilium's replacement for kube-proxy

Choose the right combination based on your application requirements, team expertise, and operational complexity tolerance. Most production environments end up using a combination of these approaches for different layers of their stack.

Securing a **Kubernetes cluster** is crucial for protecting workloads, infrastructure, and sensitive data. Here's how it can be done using **RBAC**, **Network Policies**, and **Pod Security**:

✓ 1. Role-Based Access Control (RBAC)

RBAC in Kubernetes is used to control who can perform what actions on which resources.

₩ Key Concepts:

- Roles/ClusterRoles: Define sets of permissions.
 - Role is namespace-scoped
 - ClusterRole is cluster-wide
- RoleBindings/ClusterRoleBindings: Assign roles to users or service accounts.
 - Bind users to roles within a namespace (RoleBinding)
 - Or across the cluster (ClusterRoleBinding)
- Best Practices:
 - Follow Principle of Least Privilege only give necessary access.
 - Avoid using cluster-admin role unless absolutely required.
 - Regularly audit and review permissions using:

```
kubectl auth can-i --as <user> --namespace <ns> <verb> <resource>
```

2. Network Policies

Network Policies control traffic flow between pods at the network level.

- ₩ What they do:
 - Define which pods can communicate with other pods or external endpoints.
 - Applied using labels and selectors.

Example Policy:

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
name: deny-all
```

```
namespace: my-namespace
spec:
podSelector: {} # apply to all pods
policyTypes:
- Ingress
- Egress
```

This policy denies all ingress and egress traffic.

♥ Best Practices:

- Start with a default deny-all policy and explicitly allow necessary communication.
- Apply namespace-level segmentation.
- Use labeling standards to ensure policies apply correctly.

 \triangle Note: Network Policies only work if your **CNI plugin** (like Calico or Cilium) supports them.

(PodSecurity Admission or PSA)

Kubernetes Pod Security governs what a pod is allowed to do (e.g., privilege escalation, host access).

- ₩ Pod Security Admission (PSA) Modes:
 - Enforce: Blocks pods that violate the policy.
 - Audit: Logs policy violations but does not block.
 - Warn: Warns users but allows pod creation.

◇ Pod Security Levels:

- privileged: No restrictions.
- baseline: Minimal restrictions (recommended for most apps).
- restricted: Most secure (no privilege escalation, host namespace access, etc.)

🗷 Example:

```
apiVersion: v1
kind: Namespace
metadata:
   name: secure-ns
labels:
   pod-security.kubernetes.io/enforce: restricted
   pod-security.kubernetes.io/audit: baseline
   pod-security.kubernetes.io/warn: baseline
```

PROF

● Best Practices:

- Use restricted mode for critical workloads.
- Prevent containers from:
 - Running as root (runAsNonRoot: true)
 - Accessing the host network or PID space
 - Using hostPath volumes
- Adopt SecurityContext and PodSecurityContext in pod specs.

★ Bonus Security Measures:

- Puse Secrets securely: Use tools like Sealed Secrets or Vault for managing secrets.
- **@ Enable TLS everywhere**: Encrypt API server traffic and communication between nodes.
- **[]** Audit Logs: Enable audit logging to detect unauthorized activities.

- **Regularly scan images**: Use tools like **Trivy**, **Clair**, or **Anchore** for CVE scans.
- **Use Limit access to kubelet**, etcd, and dashboard interfaces.

Terraform and Ansible are both powerful tools for infrastructure automation, but they serve different purposes and follow different **paradigms**. Here's a detailed comparison with examples to highlight how they work:

Core Differences Between Terraform and Ansible

Feature	Terraform	Ansible
Purpose	Provisioning infrastructure (laaS)	Configuration management, application setup
Language	Declarative (HCL - HashiCorp Configuration Language)	Mostly declarative, YAML (Playbooks)
State Management	Maintains state (terraform.tfstate)	Stateless (idempotent execution)
Agent Requirement	Agentless	Agentless (uses SSH or WinRM)
Best Use Cases	Cloud resource provisioning (e.g., AWS, Azure, GCP)	Installing software, patching, configuring systems
Execution Flow	Plan → Apply → Track	Task-by-task sequential execution

Terraform Example: Provision an EC2 Instance on AWS

```
# main.tf
provider "aws" {
 region = "us-east-1"
resource "aws_instance" "web_server" {
      = "ami-0c55b159cbfafe1f0"
 instance_type = "t2.micro"
 tags = {
   Name = "TerraformWeb"
 }
}
```

₩ What it does:

- Provisions an EC2 instance in AWS.
- Manages lifecycle via terraform apply / terraform destroy.

★ Ansible Example: Install Nginx on a Remote Server

```
# install-nginx.yml
- name: Install and start Nginx
 hosts: webservers
 become: yes
  tasks:
    - name: Install Nginx
     apt:
       name: nginx
        state: present
       update_cache: yes
    - name: Start Nginx
     service:
       name: nginx
       state: started
       enabled: true
```


- Connects to servers in the webservers inventory group.
- Installs and ensures Nginx is running.

Use Together for Best Results

Terraform + Ansible = Powerful Infrastructure + Configuration Management:

- 1. Use **Terraform** to:
 - Provision servers, load balancers, databases, networks, etc.
- 2. Use **Ansible** to:
 - Configure the provisioned servers (e.g., install software, apply updates).

Combined Workflow:

```
# Step 1: Provision infrastructure
terraform apply

# Step 2: Configure servers
ansible-playbook -i inventory install-nginx.yml
```

📌 Summary

Infrastructure Creation ✓ (e.g., VPCs, EC2s, RDS) ×	
Infrastructure Creation	
Configuration × (not intended)	
Cloud-native Support Strong (modular, reusable providers) Limited (uses dynamic inventories or cloud m	odules)
Orchestration Not ideal (declarative only) 🗸 Sequential task automation	
Reusability High (Modules, variables) Moderate (Roles, Playbooks)	

Implementing security best practices in DevOps pipelines is crucial for building secure, reliable, and compliant applications. This approach is often called DevSecOps, where security is integrated into every stage of the CI/CD pipeline.

Here's how you can implement key security practices focusing on:

- 1. ✓ SAST (Static Application Security Testing)
- 2. ✓ DAST (Dynamic Application Security Testing)
- 3. **⊘ Container Security**

\bigcirc 1. SAST (Static Application Security Testing)

SAST analyzes source code, bytecode, or binaries for vulnerabilities without running the code.

- Nhere it fits:
 - Early in the pipeline (pre-build or during CI)
 - As part of code reviews and merge checks
- **Practices:**
 - Integrate with IDEs or CI tools like Jenkins, GitHub Actions, GitLab CI.
 - Enforce scanning on every PR/merge request.

• Treat all issues as fail gates (e.g., fail builds on high/critical issues).

★ Common SAST Tools:

- SonarQube
- Checkmarx
- Bandit (Python)
- Semgrep

Example (GitHub Action with Semgrep):

```
- name: Run SAST using Semgrep
  uses: returntocorp/semgrep-action@v1
  with:
    config: 'p/default'
```

② 2. DAST (Dynamic Application Security Testing)

DAST tests running applications for vulnerabilities like XSS, SQL injection, and misconfigurations by simulating real attacks.

- Where it fits:
 - Post-deployment (test or staging environment)
 - After integration testing phase in CI/CD
- **Rest Practices:**
 - Run DAST scans on every build in staging.
 - Validate all APIs and user inputs.
 - Regularly scan externally exposed URLs.

★ Common DAST Tools:

- OWASP ZAP
- Burp Suite
- Nikto
- Example (ZAP CLI in CI):

docker run -t owasp/zap2docker-stable zap-baseline.py -t http://staging.example.com -r zap_report.html

© 3. Container Security

Container Security focuses on ensuring that container images are safe, non-vulnerable, and properly configured.

- Where it fits:
 - Pre-deployment and during runtime
 - Image build and registry phases
- **Rest Practices:**
 - Use minimal base images (e.g., Alpine)
 - Run containers as non-root
 - Regularly scan images for vulnerabilities
 - Sign images and enforce trust policies
 - Use read-only root filesystem and set resource limits

☆ Common Tools:

- Trivy
- Anchore
- Clair
- Docker Bench for Security
- Kube-Bench (for Kubernetes)

Example (Trivy scan):

```
trivy image myapp:latest
```

Integrating All in a DevOps Pipeline

Here's a simplified CI/CD security workflow:

```
# .github/workflows/devsecops-pipeline.yml
jobs:
 build:
   runs-on: ubuntu-latest
    steps:
     - name: Checkout code
       uses: actions/checkout@v2
      - name: Run SAST
       uses: returntocorp/semgrep-action@v1
      - name: Run Container Scan
        run:
         docker build -t myapp:latest .
         trivy image myapp:latest
      - name: Deploy to staging
        run: ./scripts/deploy.sh
      - name: Run DAST (ZAP)
        run: docker run -t owasp/zap2docker-stable zap-baseline.py -t http://staging-url -r
zap.html
```

PROF

Summary of DevSecOps Best Practices

Агеа	Key Actions
₩ SAST	Scan code for vulnerabilities early and often
⊕ DAST	Simulate real-world attacks on staging or QA environments
© Container Sec	Scan images, use signed containers, enforce least privilege
Secrets Mgmt	Use Vault, Sealed Secrets, or GitHub OIDC instead of hardcoding
■ RBAC & IAM	Implement role-based access controls across tools and infrastructure
Ш Monitoring	Use tools like Prometheus, Grafana, or ELK for security logging

Designing **Kubernetes applications for high availability (HA)** ensures that your services remain operational even if parts of the system fail. This involves both **Kubernetes cluster architecture** and **application-level strategies**.

1. Multi-Node Clusters

A multi-node cluster is foundational for high availability.

- Types of Nodes:
 - Control Plane Nodes (Masters): Manage the cluster.
 - Worker Nodes: Run your workloads (pods).

THA Strategy:

- Multiple master nodes in different zones/regions (with etcd quorum) to prevent single points of failure.
- Multiple worker nodes to distribute workloads and allow rescheduling if a node fails.

Component	HA Setup
etcd	3 or 5 nodes across zones
API server	Load-balanced across masters
Scheduler/Controller	Active-standby or HA config
Nodes	≥3 worker nodes (preferably spread across AZs)

2. Pod Replication and Distribution

Use ReplicaSets or Deployments to maintain multiple replicas of pods.

- Use replicas: 3 (or more) in Deployments.
- Use PodAntiAffinity rules to spread pods across nodes/zones:

```
affinity:
  podAntiAffinity:
  requiredDuringSchedulingIgnoredDuringExecution:
    - labelSelector:
     matchLabels:
     app: my-app
  topologyKey: "kubernetes.io/hostname"
```

• Use topologySpreadConstraints in K8s 1.18+ for even spreading across zones/nodes.

PROF

🔁 3. Service Load Balancing & Failover

Kubernetes Service Types:

- ClusterIP (default) internal only.
- NodePort / LoadBalancer external access.
- Ingress HTTP load balancing and routing.

Strategy:

- Use an external load balancer (e.g., AWS ELB, NGINX, Traefik) to:
 - Balance traffic across healthy pods.
 - Route traffic to healthy nodes.
- Leverage **readiness probes** and **liveness probes** to:
 - Prevent traffic to unhealthy pods.
 - Restart failed containers automatically.

readinessProbe: httpGet: path: /health port: 8080

initialDelaySeconds: 5
periodSeconds: 10

4. Failover Strategies

Failover ensures service continues even if something breaks.

© Key Strategies:

Layer	Strategy
Pod Level	Use Deployments or StatefulSets with multiple replicas
Node Level	Use taints/tolerations and node selectors wisely
Control Plane	Run multiple API servers behind a load balancer
Storage	Use HA volumes (e.g., AWS EBS Multi-AZ, Ceph, Portworx)
DNS	Use CoreDNS in HA mode with multiple pods

5. Disaster Recovery and Data Backup

High availability includes planning for recovery:

- Backup **etcd** regularly (etcd is the brain of the cluster).
- Use **Velero** or **Kasten K10** for volume and namespace backups.
- Design apps to be **stateless** whenever possible.

Example: Highly Available Web App Architecture on Kubernetes

- 1. **3-node control plane**, spread across zones
- 2. Auto-scaled Deployment with 3–5 replicas
- 3. NGINX Ingress controller backed by a cloud load balancer
- 4. Readiness/Liveness Probes configured
- 5. HA Persistent Volume using CSI drivers
- 6. PodAntiAffinity + topology spread constraints
- 7. **Velero** for backup and restore

***** Bonus: Key Tools for HA

Tool	Use Case
Prometheus	Monitor pod/node health
Karpenter/Cluster Autoscaler	Scale nodes automatically
Velero	Backup and disaster recovery
ExternalDNS	Manage DNS records dynamically
Istio/Linkerd	Service mesh for failover and retries

✓ Summary Checklist for Kubernetes HA:

- Multi-zone, multi-node cluster
- Multiple replicas for stateless workloads

- ✓ Load-balanced and probe-configured services
- Anti-affinity and spread policies
- ✓ HA storage and backups
- ✓ Monitoring + failover automation

Cloud-native DevOps is the practice of applying DevOps principles—automation, collaboration, continuous delivery, and monitoring—to cloud-native applications, which are designed to be scalable, resilient, and portable in cloud environments.

What is Cloud-Native DevOps?

Cloud-native DevOps combines:

Concept	Meaning
← Cloud-native	Apps built to run in cloud platforms using containers, microservices
☑ DevOps Culture and practices that automate and streamline delivery	
Automation	CI/CD, testing, infrastructure provisioning, scaling, monitoring

Together, they enable:

- Fast, frequent deployments
- Resilient, fault-tolerant systems
- Scalable, self-healing infrastructure



🚢 How Kubernetes Enables Cloud-Native DevOps

Kubernetes is the **orchestration engine** that powers many cloud-native DevOps practices:

1. *«* **Container Orchestration**

Kubernetes manages containers at scale:

- Runs and scales apps using **Pods**
- Provides **self-healing** by restarting failed containers
- Ensures high availability through replication and rolling updates



PROF

kubectl scale deployment my-app --replicas=5

2. Continuous Deployment (CD)

Kubernetes integrates seamlessly with CI/CD pipelines:

- Supports rolling updates, canary deployments, blue-green deployments
- Tools like Argo CD, Flux, or Jenkins X automate delivery to Kubernetes

strategy: type: RollingUpdate rollingUpdate: maxSurge: 1 maxUnavailable: 0

3. **☆ Infrastructure as Code (IaC)**

DevOps thrives on IaC. Kubernetes resources are YAML/JSON manifests:

- Declarative configuration
- Version-controlled infrastructure (e.g., Helm, Kustomize)

```
🔧 Example:
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
   name: app
spec:
   replicas: 3
   template:
     spec:
     containers:
        - name: app
        image: myapp:latest
```

4. 🗠 Observability and Monitoring

Kubernetes supports DevOps observability goals:

- Metrics: Prometheus + Grafana
- Logs: EFK (Elasticsearch, Fluentd, Kibana) or Loki
- Traces: Jaeger, OpenTelemetry

Q Cloud-native teams monitor:

- Pod health
- Node resource usage
- Application performance (APM)

5. Microservices Management

Kubernetes is designed to support microservice architectures:

- Each service in its own container
- Internal communication via Service Discovery
- Ingress and API Gateway manage traffic

▼ Tools like Istio/Linkerd enable:

• Traffic control

PROF

- Observability
- · Security (mTLS)

6. Self-Healing and Auto-Scaling

Kubernetes automates:

- · Rescheduling failed pods
- Horizontal and Vertical Pod Autoscaling
- Cluster Autoscaling



```
kubectl autoscale deployment web --min=3 --max=10 --cpu-percent=70
```


Cloud-Native Practice	How Kubernetes Supports It
Containerization	Manages containers, networking, and storage
Declarative IaC	YAML/Helm manifests define infrastructure
Continuous Delivery	Rolling updates, integrations with CD tools
Observability	Built-in metrics, health checks, logging
Resilience & Auto-healing	Restarts, reschedules, autoscaling
Microservices Architecture	Pod-based isolation, service mesh compatibility

Creating a complete DevOps project that integrates Docker, Kubernetes, Jenkins, Helm, and CI/CD pipelines involves building an endto-end deployment system for a cloud-native application. Here's a structured, step-by-step guide to help you set it up.



🔧 🔁 🌣 Project Overview

We'll create a pipeline that:

- 1. Builds a Docker image of a web app.
- 2. Pushes it to Docker Hub or a private registry.
- 3. Uses Helm to deploy the app to Kubernetes.
- 4. Triggers this process automatically via Jenkins CI/CD.

1. Technologies Used

Tool	Purpose
Docker	Containerize the app
Jenkins	Automate CI/CD pipelines
Kubernetes	Run the app in a containerized cluster
Helm	Package and deploy app to Kubernetes
GitHub/GitLab	Source control and webhook integration

PROF

2. Step-by-Step Implementation

Step 1: Create a Simple Web Application

Folder Structure:

```
devops-project/

    арр/

    — app.py
      requirements.txt
    L_ Dockerfile
  — helm∕
    └─ myapp/
        ├─ Chart.yaml
          — values.yaml
          - templates/
            ├─ deployment.yaml
            └─ service.yaml
```

app/app.py (Flask App):

```
from flask import Flask
app = Flask(__name__)

@app.route('/')
def hello():
    return "Hello from DevOps Pipeline!"

if __name__ == '__main__':
    app.run(host='0.0.0.0', port=5000)
```

app/requirements.txt:

```
Flask==2.0.1
```

₹ Step 2: Dockerize the App

Dockerfile:

```
FROM python:3.9-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install -r requirements.txt
COPY . .
CMD ["python", "app.py"]
```

Build & Push:

```
docker build -t your_dockerhub_username/devops-app:latest .
docker push your_dockerhub_username/devops-app:latest
```

helm/myapp/Chart.yaml:

PROF

```
apiVersion: v2
name: myapp
version: 0.1.0
```

helm/myapp/values.yaml:

```
image:
    repository: your_dockerhub_username/devops-app
    tag: latest
    pullPolicy: IfNotPresent

service:
    type: NodePort
    port: 80
    targetPort: 5000
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
 name: myapp
spec:
 replicas: 2
 selector:
   matchLabels:
     app: myapp
  template:
   metadata:
     labels:
       app: myapp
   spec:
     containers:
      - name: myapp
        image: "{{ .Values.image.repository }}:{{ .Values.image.tag }}"
        - containerPort: 5000
```

helm/myapp/templates/service.yaml:

```
apiVersion: v1
kind: Service
metadata:
   name: myapp-service
spec:
   type: {{ .Values.service.type }}
   selector:
    app: myapp
ports:
   - port: {{ .Values.service.port }}
    targetPort: {{ .Values.service.targetPort }}
```

★ Step 4: Set Up Jenkins

1. Install Jenkins Plugins:

- Docker Pipeline
- Kubernetes CLI Plugin
- Git
- Blue Ocean (optional)

2. Configure Jenkins Credentials:

- Docker Hub credentials
- Kubeconfig file for Jenkins user (as a secret file)

3. Create Jenkinsfile in root:

```
pipeline {
    agent any

environment {
        DOCKERHUB_CREDENTIALS = credentials('dockerhub-creds')
        KUBE_CONFIG = credentials('kubeconfig')
        IMAGE_NAME = "your_dockerhub_username/devops-app"
    }

stages {
    stage('Checkout') {
```

```
steps {
                git 'https://github.com/yourname/devops-project.git'
        }
        stage('Build Docker Image') {
            steps {
                script {
                    sh 'docker build -t $IMAGE_NAME:$BUILD_NUMBER ./app'
            }
        }
        stage('Push Image') {
            steps {
                script {
                    withCredentials([usernamePassword(credentialsId: 'dockerhub-creds',
usernameVariable: 'USER', passwordVariable: 'PASS')]) {
                        sh """
                        echo $PASS | docker login -u $USER --password-stdin
                        docker push $IMAGE_NAME:$BUILD_NUMBER
                    }
                }
            }
        }
        stage('Deploy to Kubernetes') {
            steps {
                script {
                    withCredentials([file(credentialsId: 'kubeconfig', variable: 'KUBECONFIG')])
{
                        sh """
                        helm upgrade --install myapp ./helm/myapp \
                          --set image.repository=$IMAGE_NAME \
                          --set image.tag=$BUILD_NUMBER \
                          --kubeconfig=$KUBECONFIG
                    }
                }
            }
        }
   }
    post {
        success {
            echo 'Deployment successful!'
        }
        failure {
            echo 'Deployment failed!'
        }
    }
}
```

PROF

Step 5: Trigger Pipeline Automatically

- Set up a **GitHub webhook** to trigger Jenkins builds on push.
- Configure polling or event-based triggers in Jenkins job settings.

Ⅲ Step 6: Monitoring and Observability

(Optional but recommended)

• Install Prometheus + Grafana in Kubernetes.

- Use tools like **Lens**, **K9s**, or **kubectl get pods** to observe deployments.
- Add **Slack notifications** in the Jenkins pipeline for alerts.


```
graph TD
A[Developer Pushes Code] --> B[Jenkins CI Triggered]
B --> C[Build Docker Image]
C --> D[Push to Docker Hub]
D --> E[Deploy with Helm to Kubernetes]
E --> F[Pods Running in Cluster]
F --> G[Users Access App via Ingress/Service]
```

Bonus: GitHub Repository Layout

```
.

— app/

— (source code + Dockerfile)

— helm/

— myapp/ (Helm chart)

— Jenkinsfile

— README.md
```

Want More?

Let me know if you'd like:

• A ready-to-fork GitHub template of this project

+ 56 / 56 **+**

- TLS/HTTPS Ingress setup
- Argo CD or GitOps version
- Monitoring & alerting setup